

← Optimization algorithms

8/10 points (80.00%)

Quiz, 10 questions

✓ Congratulations! You passed!

Next Item

1 / 1
points

1.
Which notation would you use to denote the 3rd layer's activations when the input is the 7th example from the 8th minibatch?

☐ $a^{[3]\{7\}(8)}$ ☒ $a^{[3]\{8\}(7)}$ **Correct**☐ $a^{[8]\{3\}(7)}$ ☐ $a^{[8]\{7\}(3)}$ 1 / 1
points

2.
Which of these statements about mini-batch gradient descent do you agree with?

☐ Training one epoch (one pass through the training set) using mini-batch gradient descent is faster than training one epoch using batch gradient descent.☐ You should implement mini-batch gradient descent without an explicit for-loop over different mini-batches, so that the algorithm processes all mini-batches at the same time (vectorization).☒ One iteration of mini-batch gradient descent (computing on a single mini-batch) is faster than one iteration of batch gradient descent.**Correct**1 / 1
points

3.
Why is the best mini-batch size usually not 1 and not m, but instead something in-between?

☒ If the mini-batch size is 1, you lose the benefits of vectorization across examples in the mini-batch.**Correct**

If the mini-batch size is m , you end up with stochastic gradient descent, which is usually slower than mini-batch gradient descent.

← Optimization algorithms

8/10 points (80.00%)

Un-selected is correct

☐ If the mini-batch size is 1, you end up having to process the entire training set before making any progress.

Un-selected is correct

☐ If the mini-batch size is m , you end up with batch gradient descent, which has to process the whole training set before making progress.


Correct



1 / 1
points

4.

Suppose your learning algorithm's cost J , plotted as a function of the number of iterations, looks like this:



Which of the following do you agree with?

- ☐ If you're using mini-batch gradient descent, something is wrong. But if you're using batch gradient descent, this looks acceptable.
- ☐ Whether you're using batch gradient descent or mini-batch gradient descent, something is wrong.
- ☐ Whether you're using batch gradient descent or mini-batch gradient descent, this looks acceptable.
- ☒ If you're using mini-batch gradient descent, this looks acceptable. But if you're using batch gradient descent, something is wrong.

Correct



1 / 1

points

← Optimization algorithms

8/10 points (80.00%)

5. Quiz, 10 questions

Suppose the temperature in Casablanca over the first three days of January are the same:

Jan 1st: $\theta_1 = 10^\circ C$

Jan 2nd: $\theta_2 = 10^\circ C$

(We used Fahrenheit in lecture, so will use Celsius here in honor of the metric world.)

Say you use an exponentially weighted average with $\beta = 0.5$ to track the temperature: $v_0 = 0$, $v_t = \beta v_{t-1} + (1 - \beta)\theta_t$. If v_2 is the value computed after day 2 without bias correction, and $v_2^{corrected}$ is the value you compute with bias correction. What are these values? (You might be able to do this without a calculator, but you don't actually need one. Remember what is bias correction doing.)



$$v_2 = 7.5, v_2^{corrected} = 10$$



Correct



$$v_2 = 10, v_2^{corrected} = 10$$



$$v_2 = 7.5, v_2^{corrected} = 7.5$$



$$v_2 = 10, v_2^{corrected} = 7.5$$



1 / 1

points

6.

Which of these is NOT a good learning rate decay scheme? Here, t is the epoch number.



$$\alpha = e^t \alpha_0$$



Correct



$$\alpha = \frac{1}{\sqrt{t}} \alpha_0$$



$$\alpha = 0.95^t \alpha_0$$



$$\alpha = \frac{1}{1+2^t} \alpha_0$$



0 / 1

points


7.

You use an exponentially weighted average on the London temperature dataset. You use the following to track the temperature:

$v_t = \beta v_{t-1} + (1 - \beta)\theta_t$. The red line below was computed using $\beta = 0.9$. What would happen to your red curve as you vary β ? (Check the two that apply)

8/10 points (80.00%)

Quiz, 10 questions



☐ Decreasing β will shift the red line slightly to the right.

Un-selected is correct

☐ Increasing β will shift the red line slightly to the right.

Correct

True, remember that the red line corresponds to $\beta = 0.9$. In lecture we had a green line ($\beta = 0.98$) that is slightly shifted to the right.

☐ Decreasing β will create more oscillation within the red line.

This should be selected

☐ Increasing β will create more oscillations within the red line.

This should not be selected

False. Increasing β will cause fewer oscillations



1 / 1
points


8.

Consider this figure:

← Optimization algorithms

8/10 points (80.00%)

Quiz, 10 questions



These plots were generated with gradient descent; with gradient descent with momentum ($\beta = 0.5$) and gradient descent with momentum ($\beta = 0.9$). Which curve corresponds to which algorithm?

- ☐ (1) is gradient descent with momentum (small β), (2) is gradient descent with momentum (small β), (3) is gradient descent
- ☐ (1) is gradient descent with momentum (small β), (2) is gradient descent, (3) is gradient descent with momentum (large β)
- ☐ (1) is gradient descent, (2) is gradient descent with momentum (large β), (3) is gradient descent with momentum (small β)
- ☒ (1) is gradient descent, (2) is gradient descent with momentum (small β), (3) is gradient descent with momentum (large β)

Correct

0 / 1
points

9.

Suppose batch gradient descent in a deep network is taking excessively long to find a value of the parameters that achieves a small value for the cost function $\mathcal{J}(W^{[1]}, b^{[1]}, \dots, W^{[L]}, b^{[L]})$. Which of the following techniques could help find parameter values that attain a small value for \mathcal{J} ? (Check all that apply)

- ☐ Try initializing all the weights to zero

Un-selected is correct

- ☐ Try mini-batch gradient descent

Correct

- ☐ Try better random initialization for the weights

Correct

- ☐ Try using Adam

Correct



Optimization algorithms

8/10 points (80.00%)



Quiz, 10 questions
Try tuning the learning rate α

This should be selected



1 / 1
points

10.

Which of the following statements about Adam is False?



Adam should be used with batch gradient computations, not with mini-batches.

Correct



Adam combines the advantages of RMSProp and momentum

The learning rate hyperparameter α in Adam usually needs to be tuned.We usually use "default" values for the hyperparameters β_1 , β_2 and ε in Adam ($\beta_1 = 0.9$, $\beta_2 = 0.999$, $\varepsilon = 10^{-8}$)